

Connected Components of an Undirected Graph (CUDA)

Parallel and Distributed Systems (Homework 3)

Apostolos Pantazis AM: 10910 January 2025

1 Abstract

This report details the implementation of a high-performance CUDA solver for the Connected Components (CC) problem in large-scale undirected graphs. We address the primary challenges of graph irregularity and power-law degree distributions by utilizing a Warp-Centric processing model. Algorithmic convergence is accelerated via parent-linkage and logarithmic pointer jumping. Furthermore, we eliminate the traditional I/O bottleneck by implementing an asynchronous memory-mapped (`mmap`) integer parser. Our results demonstrate sub-second convergence for the *mawi* dataset (226M vertices) and high-efficiency processing for the *com-Friendster* dataset (1.8B edges) in under 5 seconds on average.

2 Introduction

Graph processing at the scale of hundreds of millions of vertices and billions of edges requires a shift away from traditional CPU-centric traversal algorithms. Sequential methods like Depth-First Search (DFS) are inherently difficult to parallelize on GPUs due to recursion and uncoalesced memory access. Even parallel CPU frameworks like OpenCilk are often limited by the "memory wall" and synchronization overhead when dealing with irregular datasets.

This implementation leverages the NVIDIA CUDA SIMT architecture to transform the CC problem into a massively parallel label-propagation task. By mapping a single warp to each vertex row, we ensure that the high degree of "hub" nodes is handled cooperatively, maximizing memory throughput and minimizing thread divergence.

3 Methodology: GPU Parallelization

3.1 Data Structure: Compressed Sparse Row (CSR)

Contiguous memory access is maintained by representing undirected graphs in the CSR format. For undirected datasets, each edge (u, v) is stored as two directed arcs $(u \rightarrow v)$ and $(v \rightarrow u)$. This doubling of the edge list ensures that any thread analyzing vertex u can immediately identify all connected neighbors without global searches. Given that the *Friendster* dataset contains 1.8 billion edges, the CSR structure requires approximately 15 GB of VRAM, pushing the limits of modern consumer and workstation GPUs.

3.2 The Warp-Centric Model

Traditional CUDA graph kernels assign one thread to one vertex. However, in social networks, a few vertices may possess millions of neighbors while the majority have only a handful. This leads to extreme load imbalance. Our implementation assigns a **full Warp (32 threads)** to each vertex row. Threads within the warp divide the neighbor list, scanning concurrently. We utilize the `__shfl_down_sync` primitive to find the minimum label within the warp, which avoids expensive global memory atomic

operations during the scanning phase.

```

1 int warp_id = (blockIdx.x * blockDim.x + threadIdx.x) / 32;
2 int lane_id = threadIdx.x % 32;
3 if (warp_id < n) {
4     int v = warp_id;
5     int min_l = labels[v];
6     for (long long k = offsets[v] + lane_id; k < offsets[v+1]; k += 32) {
7         int neighbor_l = labels[edges[k]];
8         if (neighbor_l < min_l) min_l = neighbor_l;
9     }
10    // Intra-warp reduction using shuffle primitives
11    for (int offset = 16; offset > 0; offset /= 2) {
12        int remote = __shfl_down_sync(0xFFFFFFFF, min_l, offset);
13        if (remote < min_l) min_l = remote;
14    }
15}

```

Listing 1: Warp-Centric Label Propagation Snippet

3.3 Accelerating Convergence: Hooking and Shortcuting

Parallel label propagation is traditionally bounded by the graph diameter D . To achieve $O(\log D)$ convergence, we implemented:

- **Parent-Linkage:** Using `atomicMin` on the vertex's parent (`labels[p_v]`) effectively merges entire component trees rather than single nodes, drastically reducing the number of required kernel launches.
- **Sampling Pass:** An initial kernel samples the first two neighbors of every vertex, identifying the "Giant Component" found in most real-world networks almost instantaneously.
- **Pointer Jumping:** A compression kernel "flattens" the component trees by setting `labels[v] = labels[labels[v]]`, ensuring that every node points directly to the component root.

4 Optimized File I/O: Memory Mapping

Standard C I/O (`sscanf`, `fgets`) is the primary bottleneck for large-scale graph analytics. We redesigned the initial `readMTX` function using `mmap()` to map raw files directly into the virtual memory space. This bypasses the overhead of standard library buffering and allows for high-speed, zero-copy pointer-based parsing. For the *com-Friendster* dataset, this optimization reduced loading times from dozens of minutes to roughly 40 seconds on local SSD storage.

5 Experimental Results

The implementation was evaluated on an NVIDIA GPU environment. Presented below are the results of 10 trials for each dataset. All results utilize the optimized binary graph loader after an initial MTX parse to isolate the GPU kernel performance.

5.1 Performance Summary

Table 1: Aggregate Performance Metrics (Averaged over 10 trials)

Dataset	Vertices	Edges	Components	Avg. Kernel Time	Iterations
mawi_201512...	226,196,185	38,000,000	3,971,144	1.6634 s	4
com-Friendster	65,608,366	1,806,067,135	1	4.7923 s	3

5.2 Dataset 1: mawi_201512020330

The *mawi* dataset is characterized by a high vertex count but relatively sparse connections. The implementation achieved convergence in **four iterations**. This can be attributed to the fact that the *mawi* graph is weighted making the extra iterations necessary to be able to accurately calculate the number of strongly connected components.

Table 2: mawi_201512020330 Individual Trials

Trial	Iterations	Time (s)	Trial	Iterations	Time (s)
1	4	1.6496	6	4	1.6731
2	4	1.6203	7	4	1.6535
3	4	1.6215	8	4	1.6602
4	4	1.5984	9	4	1.7189
5	4	1.6139	10	4	1.8242
Average		1.6634			

5.3 Dataset 2: com-Friendster

The *com-Friendster* graph (1.8B edges) is the primary stress test for high-throughput label propagation. Convergence was achieved in only **three iterations**, demonstrating the power of logarithmic tree compression.

Table 3: com-Friendster Individual Trials

Trial	Iterations	Time (s)	Trial	Iterations	Time (s)
1	3	3.2530	6	3	3.7482
2	3	3.4452	7	3	5.8699
3	3	4.1613	8	3	4.4856
4	3	4.3464	9	3	5.4800
5	3	3.6154	10	3	4.6899
Average		4.3095			

6 Performance Analysis and Discussion

The experimental results confirm the effectiveness of the Warp-Centric model. The *mawi* dataset was resolved in **1.663s**, despite containing over 226 million vertices. The *com-Friendster* dataset averaged **4.79s**. We observe a significant variance in the Friendster trials (ranging from 3.6s to 5.4s), which is likely attributable to thermal throttling or PCIe bus contention during the initial transfer of the 15GB adjacency structure to the GPU memory.

Algorithmically, the convergence in 1 to 3 iterations proves that the combination of parent-linkage and pointer jumping successfully bypasses the diameter bottleneck. While standard CC implementations on Friendster often require dozens of iterations, our approach collapses the entire component structure into a single root in just 3 steps.

7 Conclusion

This implementation proves that GPU parallelization of graph algorithms can be exceptionally efficient when load imbalance and I/O bottlenecks are explicitly managed. The transition to Warp-Centric kernels ensured that even the massive hubs of social networks were processed without stalling the GPU. Coupled with memory-mapped parsing, the solution provides an end-to-end framework capable of analyzing billion-edge graphs in seconds, offering a robust tool for modern big-data analytics.

You can find all the necessary files for this project on this GitHub repository:

<https://github.com/pantazisa/Homework>